

CUDA Profiling Tutorial

March 2nd, 2020

Profiling using the gpu1-project VM

- gpu1-project VM script

```
gcloud compute instances start gpu1-project  
sudo update-alternatives --config java
```

- Mac OS or Linux:

```
gcloud compute ssh gpu1-project --ssh-flag="-Y"
```


- Windows: use nvprof and NVVP locally or on `rice.stanford.edu`

What we're doing today

- Reduction code
- Kernel 0 with NVVP
- Kernel 1 with nvprof



Reduction Kernel 0 Algorithm



Reduction Kernel 0 Code

```
template <class T>
__global__ void reduce0(T *g_idata, T *g_odata, unsigned int n)
{
    T *sdata = SharedMemory<T>();

    // load shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x * blockDim.x + threadIdx.x;

    sdata[tid] = (i < n) ? g_idata[i] : 0;

    __syncthreads();

    // do reduction in shared mem
    for (unsigned int s=1; s < blockDim.x; s *= 2)
    {
        if ((tid % (2 * s)) == 0)
            sdata[tid] += sdata[tid + s];

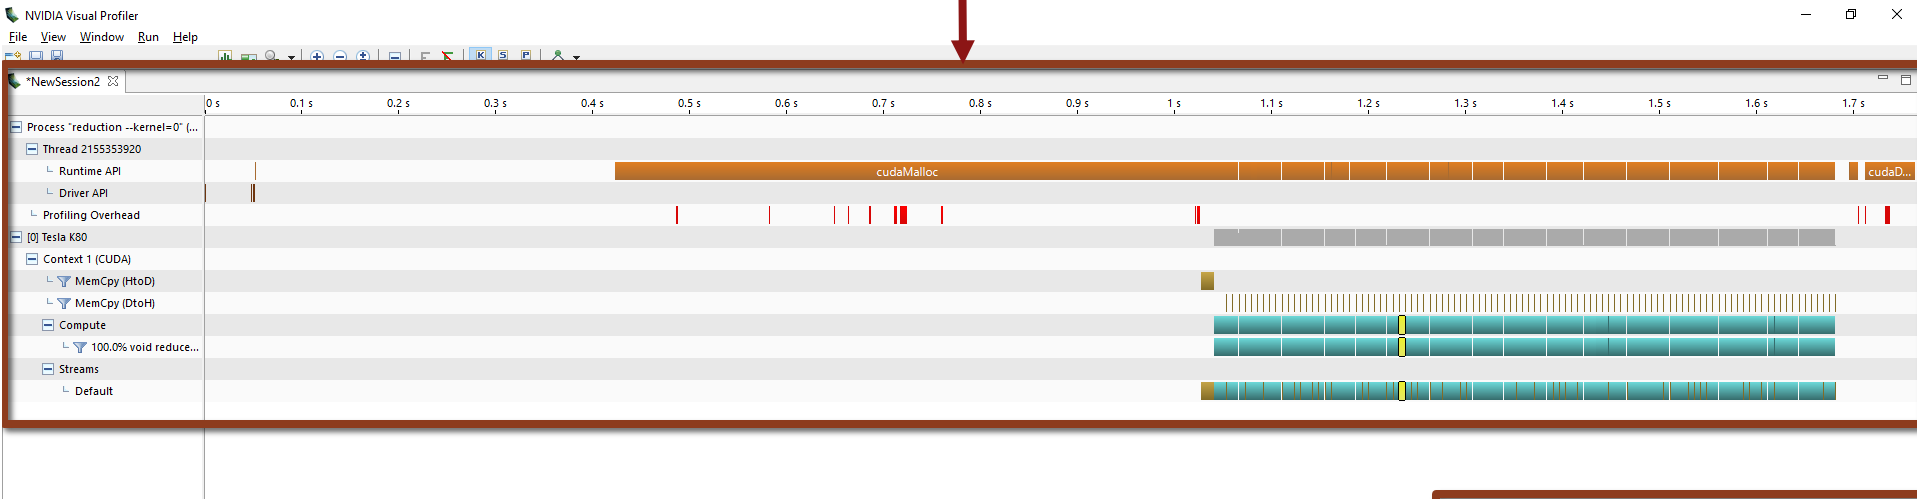
        __syncthreads();
    }

    // write result for this block to global mem
    if (tid == 0)
        g_odata[blockIdx.x] = sdata[0];
}
```

Let's see what NVVP says...

```
sudo -E /usr/local/cuda/bin/nvvp &
```

Timeline



Kernel Analysis

1. CUDA Application Analysis

The guided analysis system walks you through the various analysis stages to help you understand the optimization opportunities in your application. Once you become familiar with the optimization process, you can explore the individual analysis stages in an unguided mode. When optimizing your application it is important to fully utilize the compute and data movement capabilities of the GPU. To do this you should look at your application's overall GPU usage as well as the performance of individual kernels.

[Examine GPU Usage](#)

Determine your application's overall GPU usage. This analysis requires an application timeline, so your application will be run once to collect it if it is not already available.

[Examine Individual Kernels](#)

Determine which kernels are the most performance critical and that have the most opportunity for improvement. This analysis requires utilization data from every kernel, so your application will be run once to collect that data if it is not already available.

[Delete Existing Analysis Information](#)

If the application has changed since the last analysis then the existing analysis information may be stale and should be deleted before continuing.

[Switch to unguided analysis](#)

Timeline Detailed View

void reduce0<int>(int*, int*, unsigned int)	
Queued	n/a
Submitted	n/a
Start	1.23096 s (1,230,963,074...)
End	1.2373 s (1,237,304,501 ...)
Duration	6.34143 ms (6,341,427 ns)
Stream	Default
Grid Size	[32768, 1, 1]
Block Size	[512, 1, 1]
Registers/Thread	8
Shared Memory/Block	2 KiB
Launch Type	Normal
Occupancy	
Achieved	96.6%
Theoretical	100%
Shared Memory Configuration	
Shared Memory Executed	8 KiB
Shared Memory Bank Size	4 B

Kernel Guided Analysis

- Let's see what the profiler has to say
- Click “Examine GPU Usage”

Reduction, Throughput = 10.7144 GB/s Time = 0.00626 s, Size = 16777216 Elements, NumDevsUsed = 1, BlockSize = 512

GPU result = 2139353471

CPU result = 2139353471

TEST PASSED

⚠ **Low Kernel Concurrency** [0 ns / 682.91971 ms = 0%]

The percentage of time when two kernels are being executed in parallel is low.

[More...](#)

⚠ **Low Compute Utilization** [682.91971 ms / 458.34311 s = 0.1%]

The multiprocessors of one or more GPUs are mostly idle.

[More...](#)

i **Compute Utilization**

The device timeline shows an estimate of the amount of the total compute capacity being used by the kernels executing on the device.

What do the issues mean?

Low Memcpy/Kernel Overlap

- Means we aren't copying while running code
- Not a problem in our case

Low Kernel Concurrency

- Means we aren't executing kernels in parallel
- Not a problem in our case

Low Memcpy Overlap/Throughput

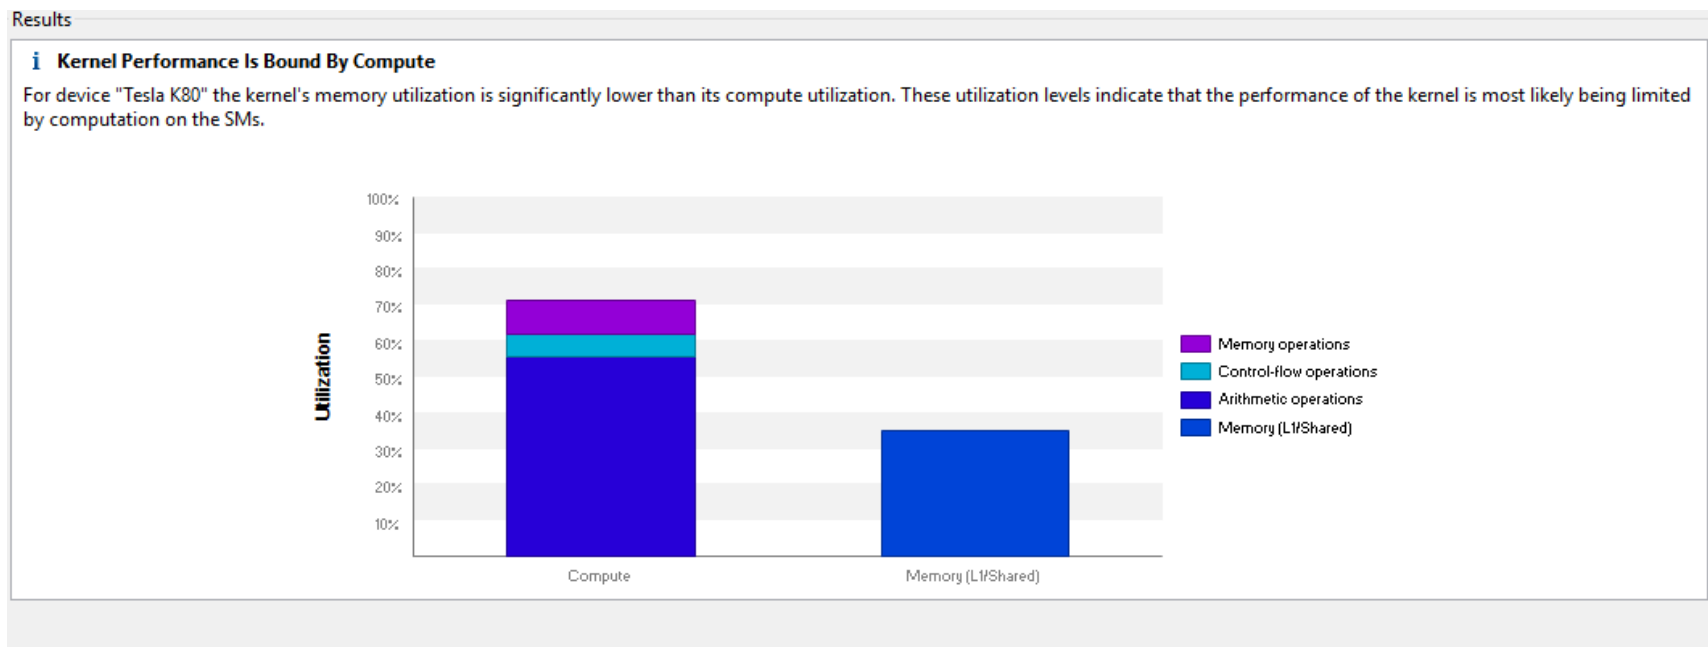
- Means we aren't copying data effectively
- Not a problem in our case

Low Compute Utilization

- Means the SMs aren't doing much work
- **This is a problem!**

What does the profiler say?

- Hit “Examine Kernels” and select the top `reduce0` kernel
- Click “Perform Kernel Analysis”
- Spending a lot of time in compute!



Tell us more, oh mighty profiler

- Hit “Perform Compute Analysis” since that’s our bottleneck

⚠ Low Warp Execution Efficiency

Warp execution efficiency is the average percentage of active threads in each executed warp. Increasing warp execution efficiency will increase utilization of the GPU's compute resources. The warp execution efficiency for these kernels is 99.9% if predicated instructions are not taken into account. The kernel's not predicated off warp execution efficiency of 72.9% is less than 100% due to divergent branches and predicated instructions.

Optimization: Reduce the amount of intra-warp divergence and predication in the kernel.

[More...](#)

⚠ Divergent Branches

Compute resource are used most efficiently when all threads in a warp have the same branching behavior. When this does not occur the branch is said to be divergent. Divergent branches lower warp execution efficiency which leads to inefficient use of the GPU's compute resources.

Optimization: Select each entry below to open the source code to a divergent branch within the kernel. For each branch reduce the amount of intra-warp divergence.

[More...](#)

▼ Line / File | NA

NA	Divergence = 6.2% [32768 divergent executions out of 524288 total executions]
----	---

i Function Unit Utilization


Different types of instructions are executed on different function units within each SM. Performance can be limited if a function unit is over-used by the instructions executed by the kernel. The following results show that the kernel's performance is not limited by overuse of any function unit.

Load/Store - Load and store instructions for local, shared, global, constant, etc. memory.

Arithmetic - All arithmetic instructions including integer and floating-point add and multiply, logical and binary operations, etc.

Control-Flow - Direct and indirect branches, jumps, and calls.

Texture - Texture operations.



Let's look at our kernel again...

```
template <class T>
__global__ void reduce0(T *g_idata, T *g_odata, unsigned int n)
{
    T *sdata = SharedMemory<T>();

    // load shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x * blockDim.x + threadIdx.x;

    sdata[tid] = (i < n) ? g_idata[i] : 0;

    __syncthreads();

    // do reduction in shared mem
    for (unsigned int s=1; s < blockDim.x; s *= 2)
    {
        if ((tid % (2 * s)) == 0)
            sdata[tid] += sdata[tid + s];


        __syncthreads();
    }

    // write result for this block to global mem
    if (tid == 0)
        g_odata[blockIdx.x] = sdata[0];
}
```

- Divergence is here
- Also, modulo is slow!

Reduction Kernel 1

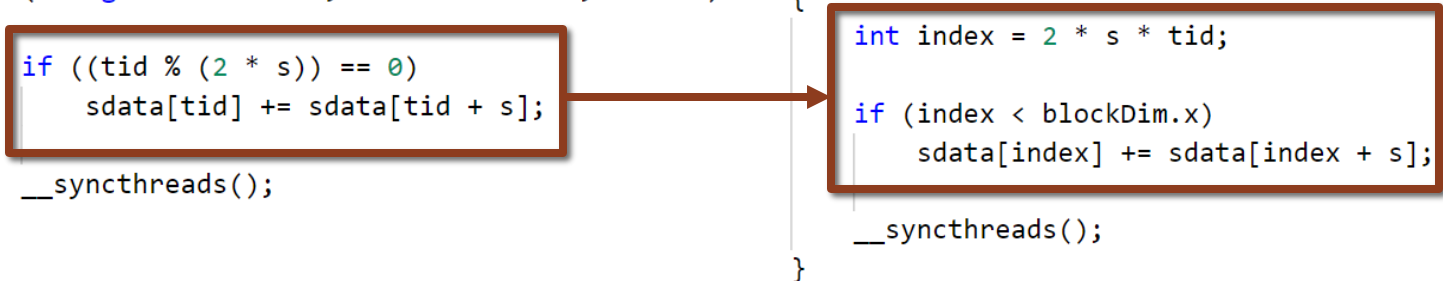
- Let's try to have only a few warps doing most of the work.
- This means less thread divergence.



Remove divergence with strided index

```
// do reduction in shared mem
for (unsigned int s = 1; s < blockDim.x; s *= 2)
{
    if ((tid % (2 * s)) == 0)
        sdata[tid] += sdata[tid + s];
    __syncthreads();
}

// do reduction in shared mem
for (unsigned int s = 1; s < blockDim.x; s *= 2)
{
    int index = 2 * s * tid;
    if (index < blockDim.x)
        sdata[index] += sdata[index + s];
    __syncthreads();
}
```



- If $s = 1$, threads 0, 2, 4, ... run
- If $s = 4$, threads 0, 8, 16 ... run

x Warp divergence

- Only consecutive threads run

✓ No divergence


Reduction Kernel 1 Profiling with nvprof

```
sudo -E /usr/local/cuda/bin/nvprof --analysis-metrics \  
-o reduction.prof ./reduction --kernel=1
```


- `--analysis-metrics` collects everything needed for NVVP Guided Analysis
- Will take some time to run

Reduction Kernel 1 Profiling Results

	Throughput GB/s
Kernel 0	14.7
Kernel 1	18.1

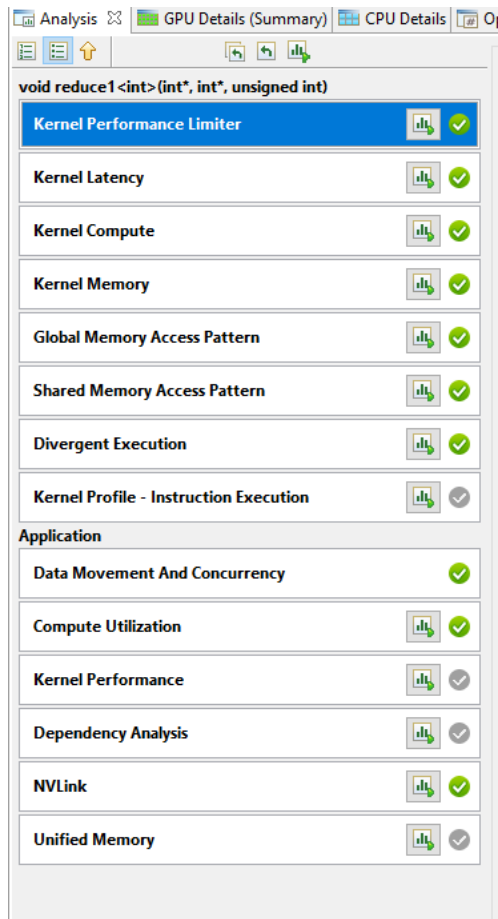


Reduction Kernel 1 Latency Analysis



Digging into Execution Dependency Stalls

- We're at the end of guided analysis
- We need more information: switch to **unguided analysis**



Different analyses based on kernel or entire application

Let's check our shared memory access pattern

107	__syncthreads();	LD.E R5, [R8];
108		I2I.S32.S32 R8, R4;
109		SHR R9, R8, 0x1f;
110	// do reduction in shared mem	MOV R8, R8;
111	for(unsigned int s=1; s < blockDim.x; s *= 2){	MOV R9, R9;
112	int index = 2 * s * tid;	MOV R10, R8;
113		MOV R11, R9;
114	if(index < blockDim.x){	MOV R10, R10;
115	sdata[index] += sdata[index + s];	MOV R11, R11;
116	}	SHF.LU64 R11, R10, 0x2, R11;
117		SHL R10, R10, 0x2;
118	__syncthreads();	IADD R10.CC, R0, R10;
119	}	IADD.X R11, R3, R11;
120		MOV R8, R10;
121	// write result for this block to global mem	MOV R9, R11;
122	if(tid == 0){	MOV R8, R8;
123	g_odata[blockIdx.x] = sdata[0];	MOV R9, R9;
124	}	LD.E R8, [R8];
125	}	IADD R5, R8, R5;
126		MOV R8, R10;
127	/*	MOV R9, R11;
128	This version uses sequential addressing -- no divergence or bank conflicts.	MOV R8, R8;
129	*/	MOV R9, R9;
130	template <class T>	ST.E [R8], R5;
131	_global_ void	NOP.S;
132	reduce2(T* g_odata, T* g_odata, unsigned int n){	

Results

Shared Memory Alignment and Access Pattern

Memory bandwidth is used most efficiently when each shared memory load and store has proper alignment and access pattern.

Optimization: Select each entry below to open the source code to a shared load or store within the kernel with an inefficient alignment or access pattern. For each load or store improve the alignment and access pattern of the memory access.

[More...](#)

Line / File reduction_kernel.cu - \home\wjen\reduction

115 Shared Load Transactions/Access = 2.8, Ideal Transactions/Access = 1 [1802240 transactions for 655360 total executions]

115 Shared Load Transactions/Access = 2.8, Ideal Transactions/Access = 1 [1802240 transactions for 655360 total executions]

115 Shared Store Transactions/Access = 2.8, Ideal Transactions/Access = 1 [1802240 transactions for 655360 total executions]

No. of banks = 8


No. of banks = 8

2 way conflict

4 way conflict


8 way conflict

8 way conflict



Reduction Algorithm, Kernel 2

- Remove bank conflicts through sequential accesses



Reduction Algorithm, Kernel 2 Implementation

```
// do reduction in shared mem
for (unsigned int s = 1; s < blockDim.x; s *= 2)
{
    int index = 2 * s * tid;

    if (index < blockDim.x)
        sdata[index] += sdata[index + s];

    __syncthreads();
}
```

```
// do reduction in shared mem
for (unsigned int s = blockDim.x / 2; s > 0; s >>= 1)
{
    if (tid < s)
        sdata[tid] += sdata[tid + s];

    __syncthreads();
}
```

**Note different loop
bounds & addition**

- 1st kernel call: we go 1/2 of block size to get next operand
- 2nd kernel call: we go 1/4 of block size to get next operand
- And so on

On your own: profile kernel 2

Thank you – any questions?